

GPU 2D Convolution Optimization Report

Anmol Gupta (27257)

November 23, 2025

1 Experimental Setup

1.1 Hardware Configuration

- GPU: NVIDIA A100-SXM4-40GB
- Memory: 40 GB HBM2
- Compute Capability: 8.0

1.2 Test Configuration

- Batch Size: 16
- Image Dimensions: 2048×2048 pixels
- Kernel Size: 11×11

2 Overall Performance Summary

Implementation	Time (ms)	Speedup	Compute Throughput (%)	Memory Throughput
Baseline	37.70	1.00×	16.81	97.88
Variant 1	7.87	4.80×	80.52	57.51
Variant 2	4.92	7.66×	89.97	49.70
Variant 3	4.75	7.94×	93.19	51.19
Variant 4	2.38	15.84×	93.03	51.09

Table 1: Overall performance comparison across all variants

3 Variant 1: Global Memory Access Pattern Optimization

3.1 Implementation Strategy

- **Approach:** Restructured thread mapping to enable memory coalescing
- **Key Design:** Aligned threadIdx.x with contiguous memory dimension (columns)
- **Block Configuration:** Changed from 16×16 to 32×8 for optimal warp alignment
- **Tradeoffs:** Pure access pattern optimization without additional memory usage

```
1 // BEFORE: Poor coalescing
2 int row = blockIdx.y * blockDim.y + threadIdx.x;
3 int col = blockIdx.x * blockDim.x + threadIdx.y;
4
5 // AFTER: Optimal coalescing
6 int col = blockIdx.x * blockDim.x + threadIdx.x; // x for columns
7 int row = blockIdx.y * blockDim.y + threadIdx.y; // y for rows
```

Listing 1: Key Optimization in Variant 1

3.2 Performance Analysis

Metric	Baseline	Variant 1
Execution Time	37.70 ms	7.87 ms
Speedup	1.00×	4.80×
Compute Throughput	16.81%	80.52%
Memory Throughput	97.88%	57.51%
Global Load Efficency	4%	1.35%

Table 2: Variant 1 performance metrics from Nsight Compute

3.3 Discussion

- **Why it works:** Enables hardware memory coalescing where 32 threads access contiguous addresses
- **Observed Phenomenon:** Dramatic compute throughput increase indicates transition from memory-bound to compute-bound
- **GPU Architecture:** Warp-based execution benefits from contiguous memory access patterns
- **Profiling Insight:** Memory throughput decreased but overall performance improved significantly

4 Variant 2: On-Chip Memory Optimization

4.1 Implementation Strategy

- **Approach:** Leveraged shared memory for data reuse and constant memory for kernel weights
- **Key Design:** Cooperative loading of input tiles with halo regions for boundary handling
- **Memory Usage:** 26×26 shared memory tiles for 11×11 kernel with 32×8 block size
- **Tradeoffs:** Increased complexity for significant data reuse benefits

```
1 __constant__ float constant_kernel[256]; // Constant memory
2 extern __shared__ float shared_tile[];    // Shared memory
3
4 // Cooperative loading with boundary handling
5 for (int idx = thread_id; idx < total_elements; idx += total_threads) {
6     if (input_row >= 0 && input_row < height && input_col >= 0 &&
7         input_col < width) {
8         shared_tile[load_y * shared_x + load_x] = pixel_val;
9     }
10}
--syncthreads();
```

Listing 2: Shared Memory Implementation in Variant 2

4.2 Performance Analysis

Metric	Variant 1	Variant 2	Change
Execution Time	7.87 ms	4.92 ms	-37.5%
Compute Throughput	80.52%	89.97%	+11.7%
Memory Throughput	57.51%	49.70%	-13.6%
Shared Memory Usage	None	26×26 tiles	New capability

Table 3: Variant 2 performance improvements

4.3 Discussion

- **Why it works:** Shared memory provides $100 \times$ higher bandwidth than global memory
- **Observed Phenomenon:** Lower memory throughput but better performance due to data reuse

- **GPU Architecture:** On-chip memories reduce pressure on global memory bandwidth
- **Profiling Insight:** Each input pixel reused $121 \times (11 \times 11 \text{ kernel})$, justifying shared memory overhead

5 Variant 3: Register-Level Optimization

5.1 Implementation Strategy

- **Approach:** Compiler-assisted optimizations with loop unrolling
- **Key Design:** Used `#pragma unroll` directives for both kernel dimensions
- **Optimization:** Enabled compiler to optimize register allocation and instruction scheduling
- **Tradeoffs:** Slight increase in register usage (32 to 34) for better ILP

```

1 // Loop unrolling for better register utilization
2 #pragma unroll
3 for (int ky = 0; ky < kernel_size; ky++) {
4     #pragma unroll
5     for (int kx = 0; kx < kernel_size; kx++) {
6         const int filter_idx = ky * kernel_size + kx;
7         const float weight = constant_kernel[filter_idx];
8         result += pixel * weight;
9     }
10 }
```

Listing 3: Register Optimization in Variant 3

5.2 Performance Analysis

Metric	Variant 2	Variant 3	Change
Execution Time	4.92 ms	4.75 ms	-3.5%
Compute Throughput	89.97%	93.19%	+3.6%
Memory Throughput	49.70%	51.19%	+3.0%
Register Usage	32	34	+6.3%

Table 4: Variant 3 performance improvements

5.3 Discussion

- **Why it works:** Eliminates loop overhead and enables better instruction-level parallelism

- **Observed Phenomenon:** Small but consistent improvement despite register pressure increase
- **GPU Architecture:** Register file provides fastest storage for thread-local computations
- **Profiling Insight:** 93.19% compute throughput indicates near-optimal compute utilization

6 Variant 4: Multi-Stream Concurrent Execution

6.1 Implementation Strategy

- **Approach:** Concurrent execution using multiple CUDA streams
- **Key Design:** Split batch into two halves processed in separate streams
- **Stream Management:** Created and synchronized two independent streams
- **Tradeoffs:** Increased complexity for significant parallelization benefits

```

1 // Use 2 streams for concurrent execution
2 const int NUM_STREAMS = 2;
3 cudaStream_t streams[NUM_STREAMS];
4
5 // Split batch and process concurrently
6 conv2d_variant3(input1, kernel, output1, batch1, height, width,
7     kernel_size, streams[0]);
8 conv2d_variant3(input2, kernel, output2, batch2, height, width,
9     kernel_size, streams[1]);
10
11 // Synchronize both streams
12 for (int i = 0; i < NUM_STREAMS; i++) {
13     cudaStreamSynchronize(streams[i]);
14     cudaStreamDestroy(streams[i]);
15 }
```

Listing 4: Multi-Stream Implementation in Variant 4

6.2 Performance Analysis

Metric	Variant 3	Variant 4	Change
Execution Time	4.75 ms	2.38 ms	-49.9%
Speedup vs Baseline	7.94×	15.84×	+7.90×
Compute Throughput	93.19%	93.03%	-0.2%
Memory Throughput	51.19%	51.09%	-0.2%

Table 5: Variant 4 performance through multi-stream execution

6.3 Discussion

- **Why it works:** Enables overlapping of independent computations across streams
- **Observed Phenomenon:** Nearly $2\times$ performance improvement while maintaining efficiency
- **GPU Architecture:** Multiple streams utilize parallel execution engines on GPU
- **Profiling Insight:** Nsight Systems timeline shows concurrent kernel execution
- **Architectural Benefit:** Better utilization of GPU compute resources through parallelism

7 Architectural Insights

7.1 Memory Hierarchy Progression

The optimization journey systematically targeted different levels of GPU memory hierarchy:

1. **Variant 1:** Global memory access patterns (coalescing)
2. **Variant 2:** Shared memory (data reuse) + Constant memory (kernel weights)
3. **Variant 3:** Register file (thread-level optimization)
4. **Variant 4:** Stream-level parallelism (concurrent execution)

7.2 Performance Bottleneck Evolution

Variant	Primary Bottleneck Addressed	Speedup
Baseline	Memory access patterns and coalescing	1.00×
Variant 1	Global memory transaction efficiency	4.80×
Variant 2	Data reuse and memory bandwidth	7.66×
Variant 3	Instruction-level parallelism	7.94×
Variant 4	Serial execution limitation	15.84×

Table 6: Evolution of performance bottlenecks through optimization

7.3 Nsight Compute Profiling Insights

Key metrics collected for each variant:

- **Compute Throughput:** Improved from 16.81% to 93.03%
- **Memory Throughput:** Evolved from inefficient high usage to efficient moderate usage
- **Memory Coalescing:** Dramatic improvement in Variant 1 maintained throughout
- **Register Usage:** Moderate increase in Variant 3 with performance benefit

8 Conclusion

8.1 Key Achievements

- **15.84× total speedup** over baseline implementation
- **93.03% compute throughput** achieved in final variant
- **Systematic optimization** through memory hierarchy and concurrency
- **Numerical correctness** maintained throughout all optimizations